

Open Data, Open Source and Open Standards in chemistry - The Blue Obelisk five years on

Jonathan Alvarsson¹, Igor V Filippov², Rajarshi Guha³, Robert Hanson⁴, Geoffrey R Hutchison⁵, Daniel M Lowe⁶, Peter Murray-Rust⁷, Noel M O'Boyle⁸, Dmitry Pavlov⁹, Ola Spjuth¹, Christoph Steinbeck¹⁰, Kevin J Theisen¹², Egon L Willighagen¹¹

¹Division of Molecular Toxicology, Institute of Environmental Medicine, Nobels vaeg 13, Karolinska Institutet, 171 77 Stockholm, Sweden

⁵Analytical and Biological Chemistry Research Facility, Cavanagh Pharmacy Building, University College Cork, College Road, Cork, Co. Cork, Ireland

⁶Cheminformatics and Metabolism Team, European Bioinformatics Institute (EBI), Wellcome Trust Genome Campus, Hinxton, Cambridge, UK

⁷Unilever Centre for Molecular Sciences Informatics, Department of Chemistry, University of Cambridge, Lensfield Road, CB2 1EW, UK

⁸GGA Software Services LLC, 41 Nab. Chernoi rechki 194342, Saint Petersburg, Russia

¹Department of Pharmaceutical Biosciences, Uppsala University, Box 591, 751 24 Uppsala, Sweden

¹⁰Department of Chemistry, University of Pittsburgh, 219 Parkman Avenue, Pittsburgh, PA 15260, USA

³NIH, Somewhere in the US, USA

¹²2iChemLabs, 200 Centennial Ave., Suite 200, Piscataway, NJ 08854, USA

Email: Jonathan Alvarsson - jonathan.alvarsson@farmbio.uu.se; Igor V Filippov - igorf@helix.nih.gov; Rajarshi Guha - guhar@mail.nih.gov; Robert Hanson - hansonr@stolaf.edu; Geoffrey R Hutchison - geoffh@pitt.edu; Daniel Lowe - dl387@cam.ac.uk; Peter Murray-Rust - pm286@cam.ac.uk; Noel M O'Boyle - n.oboyle@ucc.ie; Dmitry Pavlov - dpavlov@ggasoftware.com; ola.spjuth@farmbio.uu.se; Christoph Steinbeck - steinbeck@ebi.ac.uk; Kevin J Theisen - kevin@ichemlabs.com; Egon L Willighagen - egon.willighagen@ki.se;

* Corresponding author

Abstract

Background: The Blue Obelisk movement was established in 2005 as a response to the lack of open data, open standards and open source (ODOSOS) in chemistry. While other scientific disciplines such as physics, biology and astronomy (to name a few) were embracing new ways of doing science and reaping the benefits of community efforts, there was little if any innovation in the field of chemistry and scientific progress was actively hampered by the lack of access to data and tools.

Results: This contribution looks back on the past 5 years and surveys progress and remaining challenges in the areas of Open Source, Open Data and Open Standards in chemistry.

Conclusions: Here we show that the Blue Obelisk has been very successful in bringing together researchers and developers with common interests in ODOSOS, leading to development of many useful resources freely available to the chemistry community. But how best to engage with the wider chemistry community outside of the Blue Obelisk remains an open question.

Background

The Blue Obelisk movement was established in 2005 at the 229th National Meeting of the American Chemistry Society as a response to the lack of open data, open standards and open source (ODOSOS) in chemistry. While other scientific disciplines such as physics, biology and astronomy (to name a few) were embracing new ways of doing science and reaping the benefits of community efforts, there was little if any innovation in the field of chemistry and scientific progress was actively hampered by the lack of access to data and tools. Since 2005 it has become evident that a good amount of development in open chemical information is driven by the demands of neighbouring scientific fields. In many areas in biology, for example, the importance of small molecules and their interactions and reactions in biological systems has been realised. In fact, one of the first free and open databases and ontologies of small molecules was created as a resource about chemical structure and nomenclature by biologists. [?]

The formation of the Blue Obelisk group is somewhat unusual in that it is not a funded network, nor does it follow the industry consortium model. Rather it is a grassroots organisation, catalysed by an initial core of interested scientists, but with membership open to all who share one or more of the goals of the group:

- Open Data in Chemistry. One can obtain all scientific data in the public domain when wanted and reuse it for whatever purpose.
- Open Source in Chemistry. One can use other people's code without further permission, including changing it for one's own use and distributing it again.
- Open Standards in Chemistry. One can find visible community mechanisms for protocols and communicating information. The mechanisms for creating and maintaining these standards cover a wide spectrum of human organisations, including various degrees of consent.

Note that while some may advocate also for Open Access to publications, the Blue Obelisk goals (ODOSOS) focus more on the availability of code (to reproduce results), standards (to exchange data), and the scientific data itself. All three of these goals stem from the fundamental tenants of the scientific method for data sharing and reproducibility.

The Blue Obelisk was first described in the CDK News [?] and later as a formal paper by Guha et al. [?] in 2006. This contribution looks back on the past 5 years and surveys progress and remaining challenges in the areas of Open Data, Open Source, and Open Standards in chemistry.

Open Source

Cheminformatics toolkits

Open Source toolkits for cheminformatics have now existed for nearly ten years. During this period, some toolkits were developed from scratch in academia, whereas others were made Open Source by releasing in-house codebases under liberal licenses. When the Blue Obelisk was established five years ago, the primary toolkits under active development were the Chemistry Development Kit (CDK) [?, ?], Open Babel [?], and JOELib [?]. Of these, both the CDK and Open Babel continue to be actively developed.

[Insert main focus of work on CDK in last 5 years]

Since 2006, major new features of Open Babel include 3D structure generation and 2D structure-diagram generation, UFF and MMFF94 forcefields, and significantly expanded support for computational chemistry calculations. In addition, a major focus of Open Babel development has been to provide for accurate conversion and representation in areas of stereochemistry, kekulisation, and canonicalisation. The project has also grown, in terms of new contributors, new support from commercial companies, and second-generation tools applying Open Babel to a variety of end-user applications, from molecular editors to chemical database systems.

Two new Open Source cheminformatics toolkits have appeared since the original paper. In 2006 Rational Discovery, a cheminformatics service company (since closed down), released RDKit [?] under the BSD License. This is a C++ library with Python and (more recently) Java bindings. RDKit is actively developed and includes code donated by Novartis. Recent developments include the Java bindings, as well as performance improvements for its database cartridge. More recently, GGA Software Services (a contract programming company) released the Indigo toolkit [?] and associated software in 2009 under the GPL. Indigo is a C++ library with high-level wrappers in C, Java, Python,

and the .NET environment. Like RDKit and other toolkits, Indigo provides support for tetrahedral and cis-trans stereochemistry, 2D coordinate generation, exact/substructure/SMARTS matching, fingerprint generation, and canonical SMILES computation. It also provides some less common functionality, like matching tautomers and resonance substructures, enumeration of subgraphs, finding maximum common substructure of N input structures, and enumerating reaction products.

Second-generation tools

Although feature-rich and robust cheminformatics toolkits are useful in and of themselves, they can also be seen as providing a base layer on which additional tools and applications can be built. This is one of the reasons that cheminformatics toolkits are so important to the open source ‘ecosystem’; their availability lowers the barrier for the development of a ‘second generation’ of chemistry software that no longer needs to concern itself with the low-level details of manipulating chemical structures, and can focus on providing additional functionality and ease-of-use. Although a wide range of chemistry software has been built using Blue Obelisk components (see for example, the “Related Software” link on the Open Babel website, or “Software using CDK” at the CDK website), in this section we focus on second-generation tools which themselves have been developed by members of the Blue Obelisk.

Bioclipse [?] (v2.4 released in Aug 2010) and Avogadro [?] (v1.0 in Oct 2009) are two examples of such software, based on the CDK and Open Babel, respectively. Bioclipse is an award-winning molecular workbench for life sciences that wraps cheminformatics functionality behind user-friendly interfaces and graphical editors while Avogadro is a 3D molecular editor and viewer aimed at preparing and analysing computational chemistry calculations. Both projects are designed to be extended or scripted by users through the provision of a plugin architecture and scripting support (using Bioclipse Scripting Language [?], or Python in the case of Avogadro). An interesting aspect of both Avogadro and Bioclipse is that they share some developers with the underlying toolkits and this has driven the development of new features in the CDK and Open Babel.

Bioclipse in turn acts as base for other software such as Brunn [?], a laboratory information system for microplate based high-throughput screen-

ing. Brunn provides a graphical interface for handling different plate layouts and dilution series and can automatically generate dose response curves and calculate IC_{50} -values.

A final example of 2nd-generation Blue Obelisk software is the AMBIT2 [?] software, a GUI that facilitates registration of chemicals for the REACH EU directive on toxicity, and which is based on the CDK.

Computational chemistry analysis

Another area where the Blue Obelisk has had a significant impact in the past five years is in supporting quantum chemistry calculations and in interpreting their results. Electronic structure calculations have a long tradition in the chemistry community and a variety of programs exist, mostly proprietary software with the notable exception of NWChem. However, since each program uses different input formats, and the results are output differently even varying between different versions of the software, preparing calculations and extracting the results in an automated way is problematic.

Avogadro has already been mentioned as a GUI for preparing computational chemistry calculations. It uses Open Babel to read the output of several electronic structure packages, and also to write the input files for each. Jmol (see below) can also depict computational chemistry results including molecular orbitals.

In 2006, the Blue Obelisk project cclib [?] was established with the goal of parsing the output from computational chemistry programs and presenting it in a standard way so that further analyses could be carried out independently of the quantum package used to obtain the results. cclib is a Python library, and the current version (version 1.0.1) supports 8 different computational chemistry codes and extracts over 30 different calculated attributes. Two related BO projects build upon cclib. GaussSum [?], is a GUI that can monitor the progress of SCF and geometry convergences, and can plot predicted UV/Vis absorption and infrared spectra from appropriate logfiles containing energies and oscillator strengths for easy comparison to experimental data. QMForge [?] provides a GUI for various electronic structure analyses such as Frenking’s charge decomposition analysis [?] and Mulliken or C-squared analyses on user-defined molecular fragments. QMForge also provides a rudimentary Cartesian coordinate

editor allowing molecular structures to be saved via OpenBabel.

[Few lines on Quixote here]

Web applications

While desktop software has composed the majority of scientific tools since the computer was introduced, the internet continues to change how applications and content are distributed and presented. The web presents new opportunities for scientists as it is an open and free medium to distribute scientific knowledge, ideas and education. Web applications are software that runs within the browser, typically implemented in Java or JavaScript. Recently, a new version of the HTML specification, HTML5, defines a well-developed framework for creating native web applications in JavaScript and opens up new possibilities for visualising chemical data.

Jmol, the interactive 3D molecular viewer, is one of the most widely used chemistry applets, and indeed has seen widespread use in other fields such as biology and even mathematics (it is used for 3D depiction of mathematical functions in the Sage Mathematics Projects [?]). It is implemented in Java, and has gone from being a "Rasmol/Chime" replacement to a fully fledged molecular visualisation package, including full support for crystallography [?], display of molecular orbitals from standard basis set/coefficient data, the inclusion of dynamic minimisation using the UFF force field, and a full implementation of Daylight SMILES and SMARTS, with extensions to conformational and biomolecular substructure searching (Jmol BioSMARTS).

In 2009, iChemLabs released the ChemDoodle Web Components library [?] under the GPL v3 license (with a liberal HTML exception). This library is completely implemented in JavaScript and uses HTML5 to allow the scientist to present publication quality 2D and 3D graphics and animations for chemical structures, reactions and spectra. Beyond graphics, this tool provides a framework for user interaction to create dynamic applications through web browsers, desktop platforms and mobile devices such as the iPhone, iPad and Android devices.

The business end

Open Source provides a unique opportunity for commercial organisations to work with the cheminformatics community. Traditional business models rely

on monetisation of the source code, causing every company to repeat the development by other companies, often combined with a free (gratis) model for people working at academic institutes, to increase adoption and allow contributions from academics. This solution defines the IP on the software return on investment, but has the downside of investment losses due to duplication of software and method development, which becomes visible when proprietary companies merge.

The Blue Obelisk community, however, takes advantage that much of the investments needed for development are either paid by academic institutes and funding schemes, and by volunteers investing time and effort who get full access to the source code in return. In fact, all partners get full access up front and the Open Source licensing ensures that they will have access any time in the future. As such, it functions as a social contract between everyone to arrange the immediate return on investment. Effectively, this approach shares the burden of the high investment in having to develop cheminformatics software from scratch, allowing researchers and commercial partners alike to focus on their core business, rather than the development of prerequisites. As such, the rich collection of Open Source cheminformatics tools provided by the Blue Obelisk greatly reduces investment up front for new companies in the cheminformatics market.

The use of Open Standards allows everyone to select those Blue Obelisk components they find most useful, as they can easily replace one component with another providing the same functionality, taking advantage that they use the same standards for, for example, data exchange. This way, licensing issues are becoming a marginal problem, allowing companies to select a license appropriate for their business model. This too, allows a company to create a successful product with significantly reduced cost and effort.

At the time of writing there are many commercial companies developing chemistry solutions around Open Source cheminformatics components provided by the Blue Obelisk community. Examples of such companies include iChemLabs, IdeaConsult, Wingu, Silicos, GenettaSoft, hBar, and Inkspot Science. Some of these merely use components, but several actively contribute back to the Blue Obelisk project they use, or donate new Open Source cheminformatics projects to the community.

For example, iChemLabs released the ChemDoo-

dle Web Components library under the GPL v3 license, based on the upcoming HTML5 Open Standard. It allows making web and mobile interfaces for chemical content. The project is already being adopted by others, including iBabel by Chris Swain [?], ChemSpotlight by Geoffrey Hutchison [?] and the RSC ChemSpider [?].

Silicos has release several Open Source utilities based on OpenBabel, such as Pharaos, a tool for pharmacophore searching, Sieve for filtering molecular structure by molecular property, Stripper for removing core scaffold structures from a molecule set, and Pyramid for molecular alignment using shape determined by the Gaussian volumes as a descriptor. Additionally, contributions have been made to the OpenBabel project itself.

Other companies use Blue Obelisk components and contribute patches, smaller and larger. For example, IXELES donated the isomorphism code in the CDK, eMolecules donated canonicalisation code to OpenBabel, and AstraZeneca contributed code to the CDK for signatures. This is just a very minor selection, and the author is encouraged to contact the individual Blue Obelisk projects for an elaborate list.

Converting chemical names and images to structures

The majority of chemical information is not stored in machine-readable formats, but rather as chemical names or depictions. The OSRA and OPSIN projects focus on extracting chemical information from these sources. Such software plays a particularly important role for data mining the chemical literature, including patents and theses.

Optical Structure Recognition Application (OSRA) [?] was started in early 2007 with the goal to create the first free and open source tool for extraction and conversion of molecular images into SMILES and SD files. From the very beginning the underlying philosophy was to integrate existing open source libraries and to avoid "reinventing the wheel" wherever possible. OSRA relies on a variety of open source components: Open Babel for chemical format conversion and molecular property calculations, GraphicsMagick for image manipulation, Potrace for vectorization, GOCR and OCRAD for optical character recognition. The growing importance of image recognition technology can be seen in the fact that only a few years ago there was only

one widely available software package for chemical structure recognition - CLiDE (commercially developed at Keymodule, Ltd), but today there are as many as seven available programs.

OPSIN (Open Parser for Systematic IUPAC Nomenclature) [?] focuses instead on interpreting chemical names. The chemical name is the oldest form of communication used to describe chemicals, predating even the knowledge of the atomic structure of compounds. Chemical names are abundant in the scientific literature and encode valuable structural information. Through successive books of recommendations [?, ?], IUPAC has tried to codify and to an extent standardise naming practices. OPSIN aims to make this abundance of chemical names machine readable by translating them to SMILES, CML or InChI. The program is based around the use of a regular grammar to guide tokenisation and parsing of chemical names, followed by step-wise application of nomenclature rules. It is able to offer fast and precise conversions for the majority of names using IUPAC organic nomenclature, and is available as a web service, Java library and standalone application for maximum interoperability.

Chemical Databases

Registration, indexing and searching of chemical structures in relational databases is one of the core areas of cheminformatics. A number of structure registration systems have been published in the last five years, exploiting the fact that free cheminformatics toolkits such as OpenBabel and the CDK were available. OrChem, for example, is an extension for the Oracle 11G database that adds registration and indexing of chemical structures to support fast substructure and similarity searching. The cheminformatics functionality is provided by the Chemistry Development Kit. OrChem provides similarity searching with response times in the order of seconds for databases with millions of compounds, depending on a given similarity cut-off. For substructure searching, it can make use of multiple processor cores on today's powerful database servers to provide fast response times in equally large data sets. OrChem is free software and can be redistributed and/or modified under the terms of the GNU Lesser General Public License as published by the Free Software Foundation. All software is available via <http://orchem.sourceforge.net>.

Collaboration and interoperability

One of the effects of the Blue Obelisk has been to bring developers together from different Open Source chemistry projects so that they look for opportunities to collaborate rather than compete, and to leverage work done by other projects to avoid duplication of effort. As an example of this, when in March 2008 the Jmol development team were looking to add support for energy minimisation, rather than implement a forcefield from scratch they ported the UFF forcefield implementation from Open Babel to Jmol. This code has allowed Jmol to support 2D to 3D conversion of structures (through energy minimisation). Similarly, efficient Jmol code for atom-atom rebonding has been ported to the CDK.

Another collaborative initiative between Blue Obelisk projects was the establishment in May 2008 of the ChemiSQL project. This brought together the developers of several open source chemistry database cartridges (PgChem, MyChem, OrChem and more recently Bingo) with a view to making their database APIs more similar and collaborating on benchmark datasets for assessing performance. For two of these projects, PgChem and MyChem, which are both based on Open Babel, there is the additional possibility of working together on a shared codebase.

In the area of cheminformatics toolkits, two of the existing toolkits Open Babel and RDKit are planning to work together on a common underlying framework called MolCore [?]. This project is still in the planning stage, but if it is a success it will mean that the the two libraries will be interoperable (while retaining their existing focus) but also that the cost of maintaining the code will be shared among more developers, freeing time for the development of new features.

One of the goals of the Blue Obelisk is to promote interoperability in chemical informatics. When barriers exist to moving chemical data between different software, the community becomes fragmented and there is the danger of vendor lock-in (where users are constrained to using a particular software, a situation which puts them at a disadvantage). This applies as much to Open Source software as to proprietary software. Cinfony is a project (first release in May 2008) whose goal is to tackle this problem in the area of cheminformatics toolkits [?]. It is a Python library that enables Open Babel, the CDK, and RDKit to be used using the same API; this makes it easy, for example, to read a molecule using Open Babel, calculate descriptors using the CDK

and create a depiction using RDKit.

Another way through which interoperability of Blue Obelisk projects has been promoted and developed is through integration into workflow software such as Taverna [?] and KNIME [?] (both open source). Such software makes it easy to automate recurring tasks, and to combine analyses or data from a variety of different software and web services. A combination of the Chemistry Development Kit and Taverna, for instance, was reported in 2010 [?]. In the case of KNIME, it comes with built-in basic collection of CDK-based and Open Babel-based nodes, while other nodes for the RDKit and Indigo are available from KNIME’s “Community Updates” site.

Remaining challenges

(Say something here about benchmarks to measure accuracy. Clear examples of performance on open datasets are required. Otherwise there is nothing to counter anecdotal evidence or FUD spread by others.)

(Better engagement with industry...make it clear what and how industry members can engage with projects.)

Open Standards

InChI

The IUPAC InChI identifier is a non-proprietary and unique identifier for chemical substances designed to enable linking of diverse data compilations. Although its development predates the Blue Obelisk, software such as Open Babel has included InChI support since 2005, and support for InChI in Indigo is due in 2011.

Since the official InChI implementation is in C, it is difficult to access from the other widely used language for cheminformatics toolkits, Java. The Blue Obelisk project JNI-InChI [?] was established in 2006 to solve this problem by using the Java Native Interface to link the InChI binary to Java. In this way, it promotes the wider adoption of this standard identifier by the chemistry community.

OpenSMILES

One of the most widely used ways to store chemical structures is the SMILES format (or SMILES

string). This is a linear notation depicted by Daylight Information Systems that describes the connection table of a molecule and may optionally encode chirality. Its popularity stems from the fact that it is a compact representation of the chemical structure that is human readable and writable, and is convenient to manipulate (e.g. to include in spreadsheets, or copy from a Wikipedia article).

Despite its widespread use, a formal definition of the language did not exist beyond Daylight’s SMILES Theory Manual and tutorials. This caused some confusion in the implementation and interpretation of corner cases, for example the handling of cis/trans bond symbols at ring closures. In 2007, Craig James (eMolecules) initiated work on the OpenSMILES specification, a complete specification of the SMILES language as an Open Standard developed through a community process. The specification is largely complete and contains guidelines on reading SMILES, a formal grammar, recommendations on standard forms when writing SMILES, as well as proposed extensions.

Recently proposed CurlySMILES [?] is an extension of the SMILES notation, which allows to define crystals structures, polymers, electron delocalisation charges, molecule interactions, and many other features absent in the initial SMILES specification. There has been discussions about including parts of the CurlySMILES notation into OpenSMILES, especially polymers.

CML

?

QSAR-ML

The field of QSAR has long been hampered by the lack of open standards, which makes it difficult to share and reproduce descriptor calculations and analyses. QSAR-ML was recently proposed as an open standard for exchanging QSAR datasets [?]. A dataset in QSAR-ML includes the chemical structures (preferably described in CML) with InChI to protect integrity, chemical descriptors by linking to the Blue Obelisk Descriptor Ontology [?], response values, units, and versioned descriptor implementations to allow for integrating several descriptor software in the same calculation. Hence, a dataset described in QSAR-ML is completely reproducible. To allow for easy setup of QSAR-ML compliant

datasets, a plugin for Bioclipse was created with graphical interfaces that can be used to set up QSAR datasets and perform calculations. Descriptor implementations were initially available from CDK and JOELib, as well as via remote web services such as XMPP [?].

Remaining challenges

A core requirement for chemical structure databases and chemical registration systems in general is the notion of structure standardisation. That is, for a given input structure, multiple representations should be converted to one canonical form. Structure canonicalisation routines partially address this aspect, converting multiple alternative topologies to a single canonical form. However, the problem of standardisation is broader than just topological canonicalisation. Features that must be considered include

- topological canonicalisation
- handling of charges
- tautomer enumeration and canonicalisation
- normalisation of functional groups

Currently, most of the individual components of a ‘standardisation pipeline’ can be implemented using BO tools. The larger problem is that there is no agreed upon list of steps for a standardisation process. While some specifications have been published (e.g., PubChem) and some standardisation services and tools are available (PubChem provides an on-line service to standardise molecules and the NCGC provides a stand alone tool) each group has their own set of rules. A common reference specification for standardisation would be of immense value in interoperability between structure repositories as well as between toolkits (though the latter is still confounded by differences in lower level cheminformatic features such as aromaticity models).

We have already discussed the development of an Open SMILES standard. While much progress has been made towards a complete specification, more remains to be done before this can be considered finished. After that point, the next logical step would be to start work on a standard for the SMARTS language, the extension to SMILES that specifies patterns that match chemical substructures.

Open Data

A considerable stumbling block in advocating the release of scientific data as Open Data has been how exactly to define Open. A major step forward was the launch in 2010 of the Panton Principles for Open Data in Science [?]. This formalises the idea that Open Data maximises the possibility of reuse and repurposing, the fundamental basis of how science works. These principles recommend that published data be licensed explicitly, and preferably under CC0 (Creative Commons ‘No Rights Reserved’, also known as CCZero) [?]. This license allows others to use the data for any purpose whatsoever without any barriers. Other licenses compatible with the Panton Principles include the Open Data Commons Public Domain Dedication and Licence (PDDL), the Open Data Commons Attribution License, and the Open Data Commons Open Database License (ODbL) (see <http://www.opendefinition.org/licenses/#Data>).

Despite this positive news, little chemical data has become available from the traditional chemical fields of organic, inorganic, solid state chemistry. Table 2 lists a few notable exceptions: ChemPedia (a now discontinued crowd-sourcing project), Crystal-Eye (<http://wwmm.ch.cam.ac.uk/crystaleye/>), and the Open Notebook Science Solubility data [?]. There is also data available using licenses not compatible with the Panton Principles, but where the user is allowed to modify and redistribute the data. A new data set in this category is the data from the ChEMBL database, which is available under the Creative Commons Share-Alike Attribution license [?].

Importantly, publishing data as CC0 is becoming easier now that websites are becoming available to simplify publishing data. Two projects that can be mentioned in this context are FigShare (<http://figshare.com/>), where the data behind unpublished figures can be hosted, and Dryad (<http://datadryad.org/>) where data behind publications can be hosted. Initiatives like this make it possible to host small amounts of data, and those combined are expected to become soon a substantial knowledge base.

Other areas of activity

While each Blue Obelisk project has its own website and point of contact (typically a mailing list), because of the breadth of BO projects it can be difficult for a newcomer to understand which of

them, if any, can best address a particular problem. To address this issue, members of the Blue Obelisk established a Question & Answer website at <http://blueobelisk.shapado.com> (see Figure 3). This is a website in the style of Stack Overflow that encourages high quality answers (and questions) through the use of a voting system. In the year since it was established, over 200 users have registered, many of whom had no previous involvement with the Blue Obelisk, showing that the Q&A website complements earlier existing channels of communication.

The rise of self-publishing and print-on-demand services has meant that publishing a book is now as straightforward as uploading to an appropriate website. Unlike the traditional publishing route where books with projected low sales volume would be expensive, websites such as Lulu allow the sale of low-priced books on chemistry software, and books are now available for purchase on Jmol [?], the Chemistry Development Kit [?] and Open Babel [?].

Conclusions

We have shown that the Blue Obelisk has been very successful in bringing together researchers and developers with common interests in ODOSS, leading to development of many useful resources freely available to the chemistry community. Figure 2 shows how the various Blue Obelisk projects collaborate. But how best to engage with the wider chemistry community outside of the Blue Obelisk remains an open question. If the Blue Obelisk is truly to make an impact, then an attempt must be made to reach beyond the subscribers to the BO mailing list and blogs of members.

We hope to see this involvement between the Blue Obelisk and the wider community grow in the future. To this end, we encourage the reader to visit <http://blueobelisk.org>, send a message to our mailing list, investigate related projects or read our blogs.

Authors contributions

Charles Darwin did all the work. The others stole the glory.

Acknowledgements

Thanks to everyone.

Figures

Figure 1 - Blue Obelisk logo

Figure 2 - Dependency diagram of Blue Obelisk projects.

Each block represents a project. Square blocks show Open Data, ovals are Open Source, and diamonds are Open Standards. Colors represent license: LGPL is green, GPL is orange, and BSD is blue.

Figure 3 - Screenshot of the Blue Obelisk eXchange Question and Answer website.

Tables

Table 1 - Blue Obelisk Open Source software projects

(Description if necessary XXXXXXXXXXXXXXXX. Add citations to project names.)

Name	Website	Des
Cheminformatics toolkits		
Chemistry Development Kit (CDK) [?, ?]	http://cdk.sf.net	
Cinfony	http://cinfony.googlecode.com	Noel O'
Indigo	http://ggasoftware.com/opensource/indigo	
Open Babel	http://openbabel.org	
RDKit	http://rdkit.org	
ChemDoodle Web Components	http://web.chemdoodle.com	
Integration		
CDK-Taverna [?]	http://cdk-taverna.blah.XXX	CH
A3	..	
Interconversion		
OSRA	http://osra.sf.net	Igor
OPSIN	http://opsin.ch.cam.ac.uk	Da
A3	..	
Structure Databases		
OrChem [?]	http://orchem.sourceforge.net	Christoph Steinbe
MyChem	http://XXX.sourceforge.net	Jerome Pansanel
PGChem	http://XXX.sourceforge.net	Ernst-Georg Schmid
Bingo	http://XXX.XXX.XXX	Dmitry Pavlov,
cclib	http://cclib.sourceforge.net	Python library for parsin

Table 2 - Open Data in chemistry.

Overview of major open chemical data available under a license or waiver compatible with the Panton Principles.

Name	License/Waiver	Description
ChemPedia	CC0	Crowd-sourced chemical names. Project discontinued.
CrystalEye	PPDL	Crystal structures from primary literature.
ONS Solubility	CC0	Solubility data for various solvents.

Additional Files

Additional file 1 — Sample additional file title

Additional file descriptions text (including details of how to view the file, if it is in a non-standard format or the file extension). This might refer to a multi-page table or a figure.

Additional file 2 — Sample additional file title

Additional file descriptions text.